

Machine Learning and Pattern Recognition


A High Level Overview

Prof. Anderson Rocha


(Main bulk of slides kindly provided by **Prof. Sandra Avila**
and based on materials by Fei-Fei Li & Justin Johnson & Serena Yeung)
Institute of Computing (IC/Unicamp)

Convolutional Neural Networks (CNNs)

Fully Connected Layer




CIFAR-10




$32 \times 32 \times 3$ image \Rightarrow stretch to 3072×1


Fully Connected Layer




CIFAR-10



$32 \times 32 \times 3$ image \Rightarrow stretch to 3072×1





What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter


What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter




What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4		

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &0*1 + 0*0 + 1*1 = 4 \end{aligned}$$


What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	

$$\begin{aligned} &1*1 + 1*0 + 0*1 + \\ &1*0 + 1*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 3 \end{aligned}$$


What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4

$$\begin{aligned} &1*1 + 0*0 + 0*1 + \\ &1*0 + 1*1 + 0*0 + \\ &1*1 + 1*0 + 1*1 = 4 \end{aligned}$$

What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4
2		

$$\begin{aligned} &0*1 + 1*0 + 1*1 + \\ &0*0 + 0*1 + 1*0 + \\ &0*1 + 0*0 + 1*1 = 2 \end{aligned}$$

What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4
2	4	

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 4 \end{aligned}$$


What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4
2	4	3

$$\begin{aligned} &1*1 + 1*0 + 0*1 + \\ &1*0 + 1*1 + 1*0 + \\ &1*1 + 1*0 + 0*1 = 3 \end{aligned}$$

What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4
2	4	3
2		

$$\begin{aligned} &0*1 + 0*0 + 1*1 + \\ &0*0 + 0*1 + 1*0 + \\ &0*1 + 1*0 + 1*1 = 2 \end{aligned}$$

What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1

3×3 filter



4	3	4
2	4	3
2	3	

$$\begin{aligned} &0*1 + 1*0 + 1*1 + \\ &0*0 + 1*1 + 1*0 + \\ &1*1 + 1*0 + 0*1 = 3 \end{aligned}$$

What is a Convolution?

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

5×5 matrix

1	0	1
0	1	0
1	0	1


3×3 filter




4	3	4
2	4	3
2	3	4

$$\begin{aligned} &1*1 + 1*0 + 1*1 + \\ &1*0 + 1*1 + 0*0 + \\ &1*1 + 0*0 + 0*1 = 4 \end{aligned}$$

What is a Convolution?



0	1	0
1	-4	1
0	1	0



What is a Convolution?



Edge
Detection

0	1	0
1	-4	1
0	1	0




What is a Convolution?




Emboss

-2	-1	0
-1	1	1
0	1	2




Convolution Layer

$32 \times 32 \times 3$ image \Rightarrow preserve spatial structure



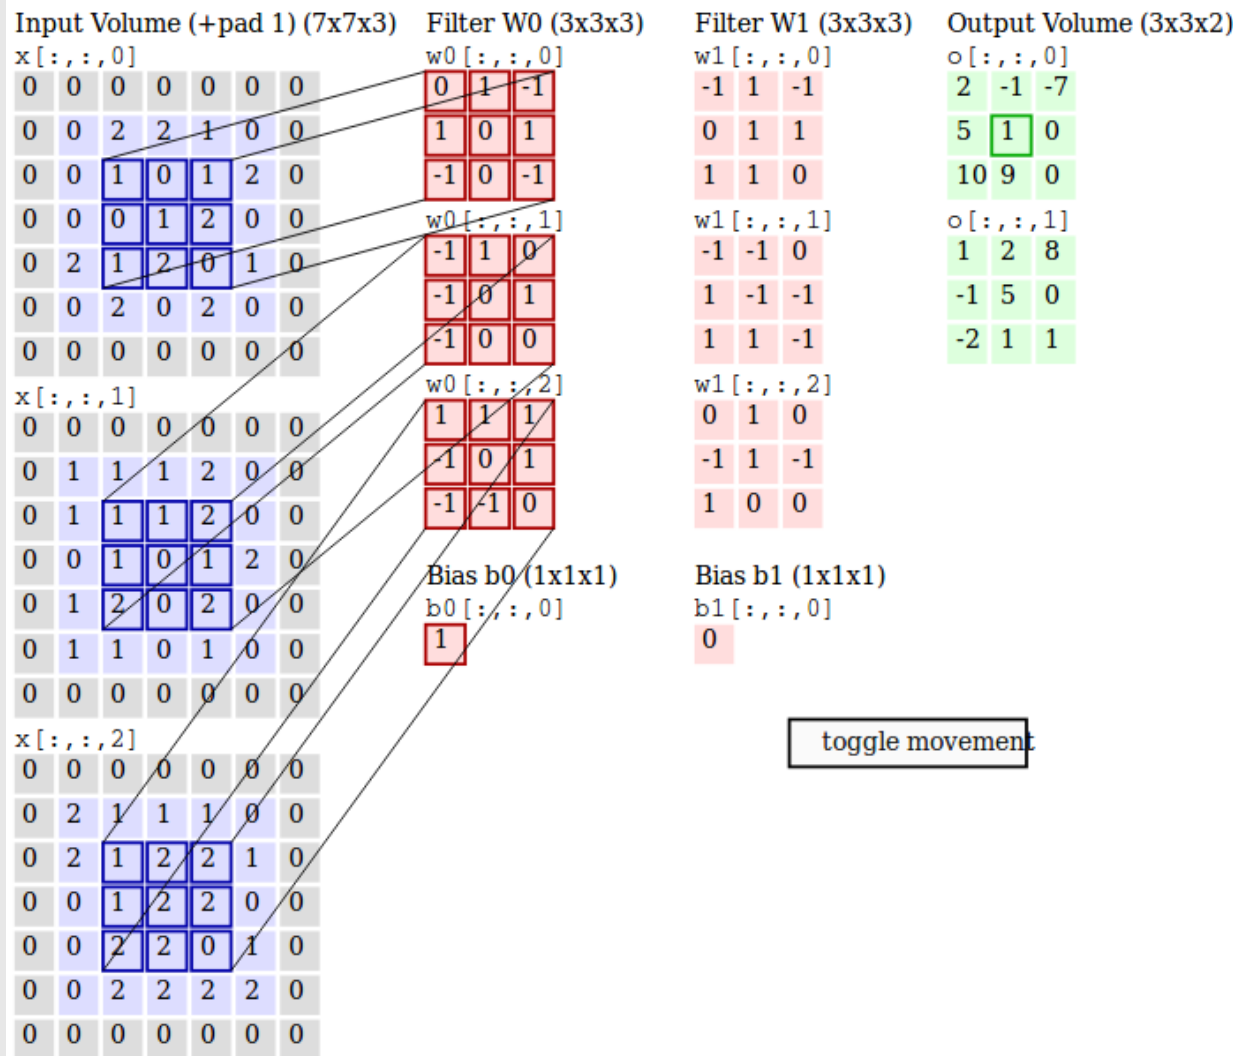
Convolve the filter with the image i.e. “slide over the image spatially, computing dot products”

$5 \times 5 \times 3$ filter \rightarrow




Filters always extend the full depth of the input volume

<http://cs231n.github.io/convolutional-networks>




Convolution Layer

$32 \times 32 \times 3$ image \Rightarrow preserve spatial structure




Convolve (slide) over all
spatial locations



$32 \times 32 \times 3$ image
 $5 \times 5 \times 3$ filter

If we had 6 $5 \times 5 \times 3$ filters ...

6 activation maps



Convolutional Layer

The size of the **Activation Map** (or Feature Map or Convolved Feature) is controlled by three parameters:

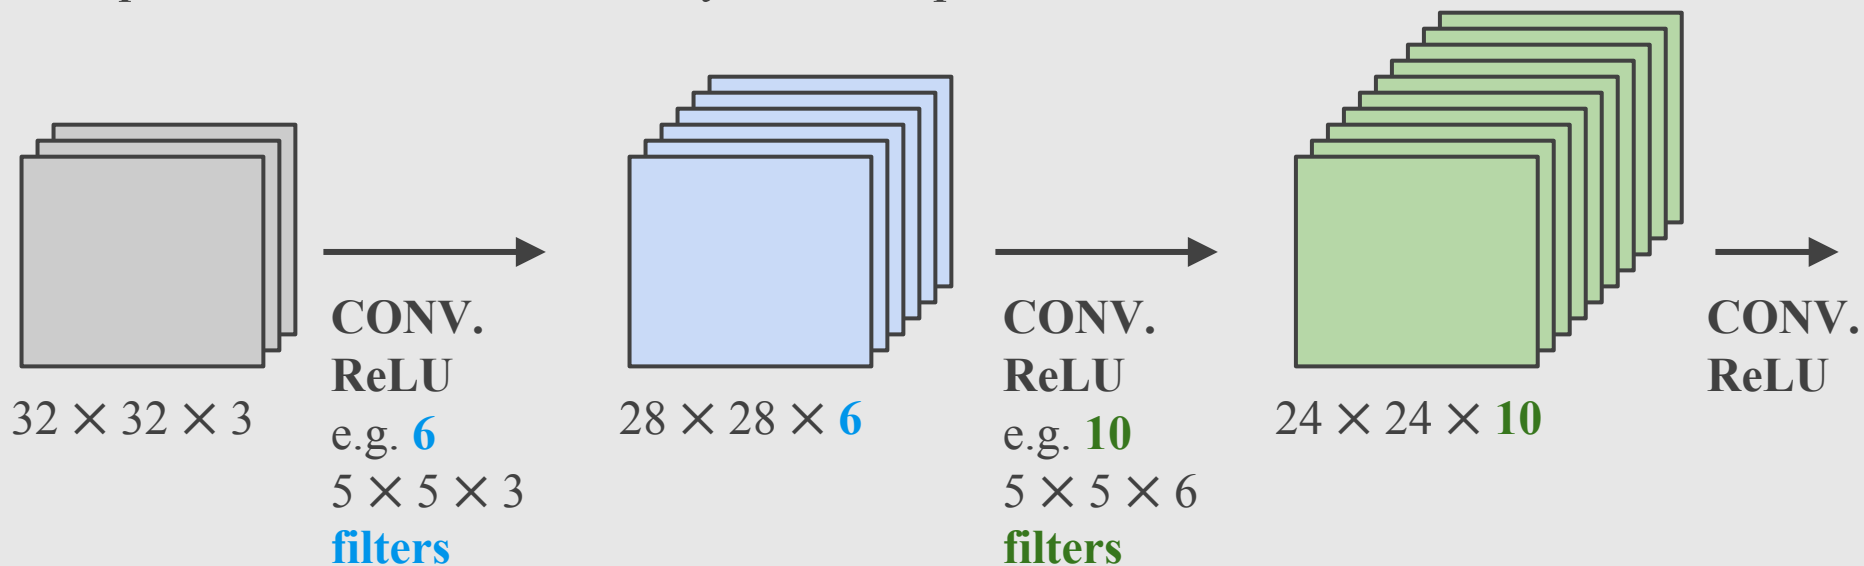
Convolutional Layer

The size of the **Activation Map** (or Feature Map or Convolved Feature) is controlled by three parameters:

- **Depth**: corresponds to the **number of filters** we use for the convolution operation.

Convolutional Layers

Sequence of Convolutional Layers, interspersed with activation functions.




Convolutional Layer

The size of the **Activation Map** (or Feature Map or Convolved Feature) is controlled by three parameters:


- **Depth**: corresponds to the **number of filters** we use for the convolution operation.
- **Stride**: the **number of pixels by which we slide** our filter matrix over the input matrix.

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**

Convolutional Layer: Stride



7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**


Convolutional Layer: Stride



7×7 input (spatially)
assume 3×3 filter
applied with **stride 1**


\Rightarrow **5×5 output**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 2**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 2**

Convolutional Layer: Stride




7×7 input (spatially)
assume 3×3 filter
applied with **stride 2**

Convolutional Layer: Stride



7×7 input (spatially)
assume 3×3 filter
applied with **stride 2**

Convolutional Layer: Stride




7×7 input (spatially)

assume 3×3 filter

applied with **stride 2**

\Rightarrow **3×3 output**


Convolutional Layer: Stride



Output size:

$$(N - F) / \text{stride} + 1$$

Convolutional Layer: Stride



Output size:

$$(N - F) / \text{stride} + 1$$

e.g. $N = 7, F = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3) / 1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3) / 2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3) / 3 + 1 = 2.33$$

Convolutional Layer

The size of the **Activation Map** (or Feature Map or Convolved Feature) is controlled by three parameters:

- **Depth**: corresponds to the **number of filters** we use for the convolution operation.
- **Stride**: the **number of pixels by which we slide** our filter matrix over the input matrix.
- **Zero-padding**: sometimes, it is convenient to pad the input matrix with **zeros around the border**.

Convolutional Layer: Zero-Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

7×7 input,
 3×3 filter applied
with **stride 1 with pad 1**

What is the output?
 7×7 output

Convolutional Layer: Zero-Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

In general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$ **(will preserve size spatially)**.


e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Shrinking too fast is not good, doesn't work well.

$32 \rightarrow 28 \rightarrow 24 \rightarrow \dots$



Number of Parameters

Input volume: $32 \times 32 \times 3$

10 5×5 filters with stride 1, pad 2

Number of parameters in this layer?

Number of Parameters


Input volume: $32 \times 32 \times 3$

$10 \times 5 \times 5$ filters with stride 1, pad 2

Number of parameters in this layer?


Each filter has $5 \times 5 \times 3 + 1 = 76$ parameters (+1 for bias)

$$\Rightarrow 76 \times 10 = 760$$




Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently



Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently




Pooling Layer

The function of Pooling is **to progressively reduce the spatial size of the input representation**. In particular, pooling

- makes the input representations (feature dimension) smaller and more manageable
- reduces the number of parameters and computations in the network, therefore, controlling overfitting
- makes the network invariant to small transformations, distortions and translations in the input image

Fully Connected Layer

- Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



<http://scs.ryerson.ca/~aharley/vis/conv/flat.html>

Draw your number here



Downsampled drawing: **3**

First guess: **3**

Second guess: **7**

Layer visibility

Input layer

Show

Convolution layer 1

Show



Visualizing a CNN trained on Handwritten Digits


- Input image: 1024 pixels (32 x 32 image)
- CONV 1 (+ RELU): 6 5×5 (stride 1) filters
- POOL 1: 2×2 max pooling (with stride 2)
- CONV 2 (+ RELU): 16 5×5 (stride 1) filters
- POOL 2: 2×2 max pooling (with stride 2)
- 3 FC layers:
 - 120 neurons in the first FC layer
 - 100 neurons in the second FC layer
 - 10 neurons in the third FC (Output layer)

DNNs Architectures


DNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)

ImageNet 2012 — AlexNet, 8 layers



ImageNet 2014 — GoogLeNet, 22 layers




ImageNet 2015 — ResNet, 152 layers



DNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)

LeNet-5 [LeCun et al., 1998]




Convolution filters: 5x5 with stride 1

Subsampling (Pooling) layers: 2x2 with stride 2

[CONV-POOL-CONV-POOL-FC-FC]

AlexNet [Krizhevsky et al., 2012]



AlexNet [Krizhevsky et al., 2012]

Architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4


CONV5

MAX POOL3

FC6

FC7

FC8



AlexNet [Krizhevsky et al., 2012]

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1

[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1


[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

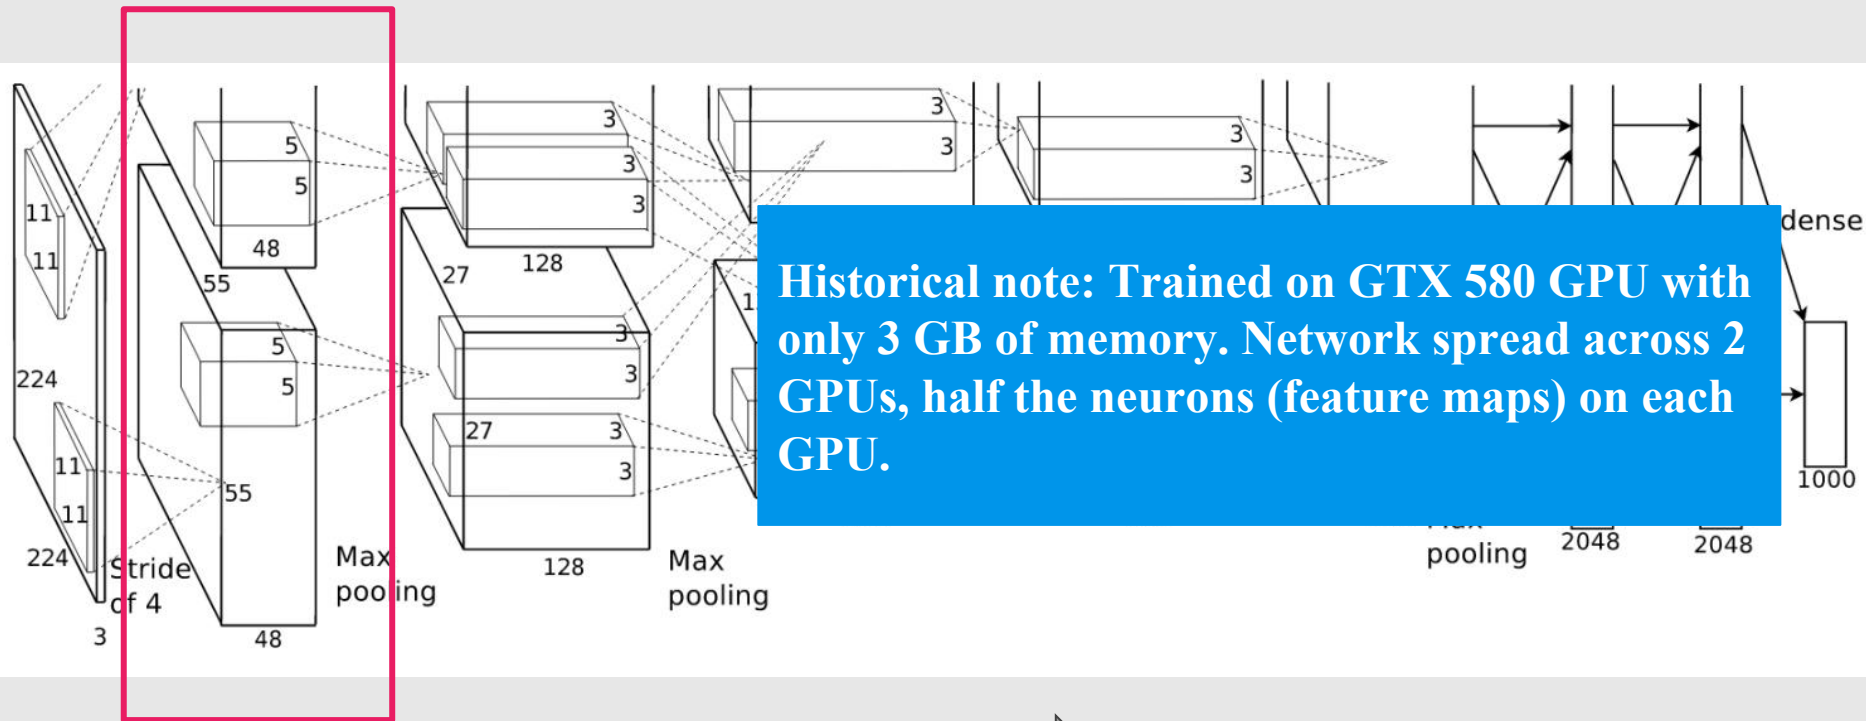
[1000] **FC8**: 1000 neurons (class scores)

AlexNet [Krizhevsky et al., 2012]



$[55 \times 55 \times 96]$ **CONV1** $\rightarrow [55 \times 55 \times 48] \times 2$

AlexNet [Krizhevsky et al., 2012]



$[55 \times 55 \times 96]$ CONV1 $\rightarrow [55 \times 55 \times 48] \times 2$

AlexNet [Krizhevsky et al., 2012]


Details:

- 60 million learned parameters
- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- 7 CNN ensemble: 18.2% \rightarrow 15.4%


Dropout

- Dropout is a radically different technique for regularization.
- Dropout doesn't rely on modifying the cost function. Instead, **in dropout we modify the network itself.**

Dropout



Standard Network



After applying dropout

Dropout

- Heuristically, when we dropout different sets of neurons, it's rather like we're training different neural networks.
- The dropout procedure is like averaging the effects of a very large number of different networks.
- The different networks will overfit in different ways, and so, hopefully, the net effect of dropout will be to reduce overfitting.

AlexNet [Krizhevsky et al., 2012]

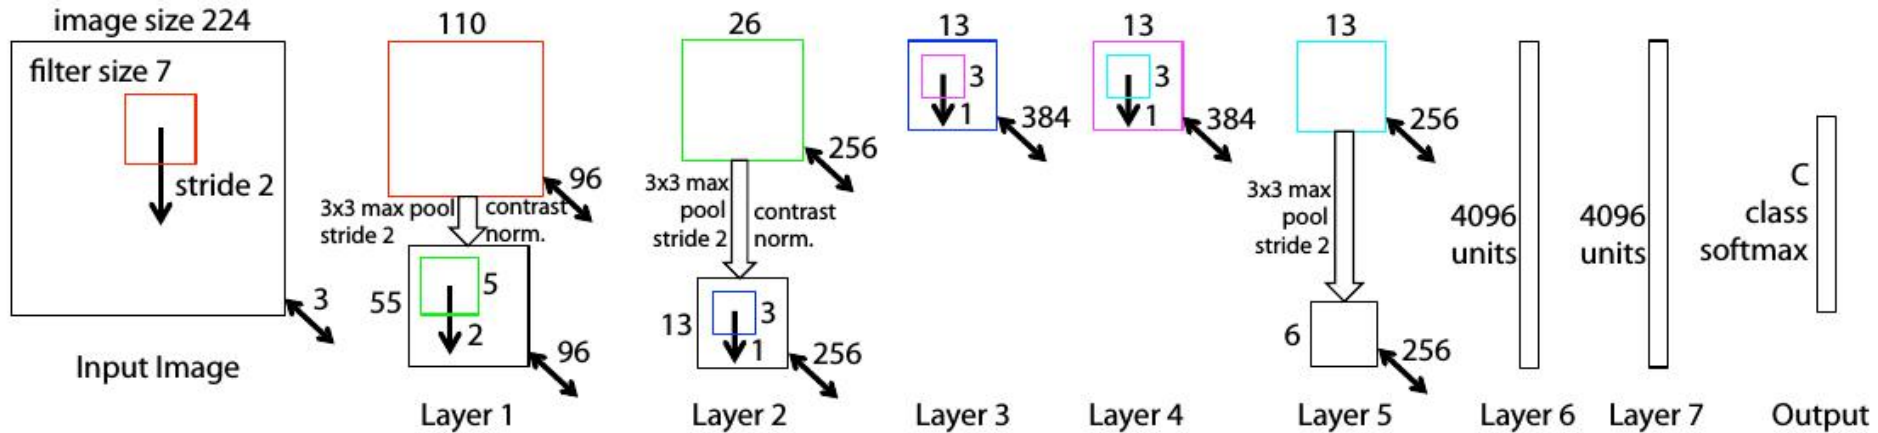
Details:

- 60 million learned parameters
- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- 7 CNN ensemble: 18.2% \rightarrow 15.4%

DNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)

ZFNet [Zeiler & Fergus, 2013]



AlexNet but:

CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

DNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)

VGGNet [Simonyan & Zisserman, 2014]

Small filters, Deeper networks


8 layers (AlexNet)

16 - 19 layers (VGG16Net)


Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

11.7% in ILSVRC'13 (ZFNet)


7.3% in ILSVRC'14



AlexNet



VGG16



VGG19

DNNs Architectures

- **LeNet** by Yann LeCun, Léon Bottou & Yoshua Bengio (1998)
- **AlexNet** by Alex Krizhevsky, Ilya Sutskever & Geoff Hinton (2012)
- **ZF Net** by Matthew Zeiler & Rob Fergus (2013)
- **GoogLeNet** by Szegedy et al. (2014)
- **VGGNet** by Karen Simonyan & Andrew Zisserman (2014)
- **ResNet** by Kaiming He et al. (2015)

To be continued ...

References

— — —

Machine Learning Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chap. 11 & 13

Machine Learning Courses

- <https://www.coursera.org/learn/neural-networks>
- “The 3 popular courses on Deep Learning”: <https://medium.com/towards-data-science/the-3-popular-courses-for-deeplearning-ai-ac37d4433bd>